

# v-milesstone-heart-test-2

September 23, 2024

## 1 MACHINE LEARNING

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import \
    mean_squared_error, r2_score, confusion_matrix, classification_report, accuracy_score
from sklearn.preprocessing import StandardScaler
```

### 1.1 Exploring the CAR Dataset with different (REGRESSION ALGORITHMS)

We have a data which is car or not according to features in it. We will try to use this data to create a model which tries predict types of car and price. We will use regression algorithms.

```
[2]: df = pd.read_csv("C:/kgisl class/MILESTONE - 2/heart.csv")
df
```

```
[2]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
0	52	1	0	125	212	0	1	168	0	1.0	
1	53	1	0	140	203	1	0	155	1	3.1	
2	70	1	0	145	174	0	1	125	1	2.6	
3	61	1	0	148	203	0	1	161	0	0.0	
4	62	0	0	138	294	1	1	106	0	1.9	
...	...	...	...	...	...	...	...	...	...	...	...
1020	59	1	1	140	221	0	1	164	1	0.0	
1021	60	1	0	125	258	0	0	141	1	2.8	
1022	47	1	0	110	275	0	0	118	1	1.0	
1023	50	0	0	110	254	0	0	159	0	0.0	
1024	54	1	0	120	188	0	1	113	0	1.4	

	slope	ca	thal	target
0	2	2	3	0
1	0	0	3	0
2	0	0	3	0

3	2	1	3	0
4	1	3	2	0
...	...	...	...	...
1020	2	0	2	1
1021	1	1	3	0
1022	1	1	2	0
1023	2	0	2	1
1024	1	1	3	0

[1025 rows x 14 columns]

## Exploring the dataset

```
[3]: df.shape
```

```
[3]: (1025, 14)
```

```
[4]: df.dtypes
```

```
[4]: age          int64
sex            int64
cp             int64
trestbps       int64
chol           int64
fbs            int64
restecg        int64
thalach        int64
exang          int64
oldpeak        float64
slope          int64
ca             int64
thal           int64
target         int64
dtype: object
```

```
[5]: df.head()
```

```
[5]:   age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  slope  \
0   52   1   0     125    212   0         1     168     0       1.0     2
1   53   1   0     140    203   1         0     155     1       3.1     0
2   70   1   0     145    174   0         1     125     1       2.6     0
3   61   1   0     148    203   0         1     161     0       0.0     2
4   62   0   0     138    294   1         1     106     0       1.9     1

   ca  thal  target
0   2     3       0
1   0     3       0
```

2	0	3	0
3	1	3	0
4	3	2	0

```
[6]: df.tail()
```

```
[6]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
1020	59	1	1	140	221	0	1	164	1	0.0	
1021	60	1	0	125	258	0	0	141	1	2.8	
1022	47	1	0	110	275	0	0	118	1	1.0	
1023	50	0	0	110	254	0	0	159	0	0.0	
1024	54	1	0	120	188	0	1	113	0	1.4	

	slope	ca	thal	target
1020	2	0	2	1
1021	1	1	3	0
1022	1	1	2	0
1023	2	0	2	1
1024	1	1	3	0

```
[7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1025 non-null   int64
1   sex         1025 non-null   int64
2   cp          1025 non-null   int64
3   trestbps    1025 non-null   int64
4   chol        1025 non-null   int64
5   fbs         1025 non-null   int64
6   restecg     1025 non-null   int64
7   thalach     1025 non-null   int64
8   exang       1025 non-null   int64
9   oldpeak     1025 non-null   float64
10  slope       1025 non-null   int64
11  ca          1025 non-null   int64
12  thal        1025 non-null   int64
13  target      1025 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
```

```
[8]: df.describe()
```

```
[8]:
```

	age	sex	cp	trestbps	chol \
count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000
mean	54.434146	0.695610	0.942439	131.611707	246.000000
std	9.072290	0.460373	1.029641	17.516718	51.59251
min	29.000000	0.000000	0.000000	94.000000	126.000000
25%	48.000000	0.000000	0.000000	120.000000	211.000000
50%	56.000000	1.000000	1.000000	130.000000	240.000000
75%	61.000000	1.000000	2.000000	140.000000	275.000000
max	77.000000	1.000000	3.000000	200.000000	564.000000

	fbs	restecg	thalach	exang	oldpeak \
count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000
mean	0.149268	0.529756	149.114146	0.336585	1.071512
std	0.356527	0.527878	23.005724	0.472772	1.175053
min	0.000000	0.000000	71.000000	0.000000	0.000000
25%	0.000000	0.000000	132.000000	0.000000	0.000000
50%	0.000000	1.000000	152.000000	0.000000	0.800000
75%	0.000000	1.000000	166.000000	1.000000	1.800000
max	1.000000	2.000000	202.000000	1.000000	6.200000

	slope	ca	thal	target
count	1025.000000	1025.000000	1025.000000	1025.000000
mean	1.385366	0.754146	2.323902	0.513171
std	0.617755	1.030798	0.620660	0.500070
min	0.000000	0.000000	0.000000	0.000000
25%	1.000000	0.000000	2.000000	0.000000
50%	1.000000	0.000000	2.000000	1.000000
75%	2.000000	1.000000	3.000000	1.000000
max	2.000000	4.000000	3.000000	1.000000

```
[9]: df.columns
```

```
[9]: Index(['age', 'sex', 'cp', 'trestbps', 'chol', 'fbs', 'restecg', 'thalach',
          'exang', 'oldpeak', 'slope', 'ca', 'thal', 'target'],
          dtype='object')
```

```
[10]: df.isnull().sum()
```

```
[10]: age      0
      sex      0
      cp       0
      trestbps  0
      chol     0
      fbs      0
      restecg   0
      thalach   0
      exang     0
```

```

oldpeak      0
slope        0
ca           0
thal         0
target       0
dtype: int64

```

### 1.1.1 Exploratory data analysis - (EDA)

```
[11]: df.corr()
```

```

[11]:
      age      sex      cp  trestbps      chol      fbs  \
age      1.000000 -0.103240 -0.071966  0.271121  0.219823  0.121243
sex     -0.103240  1.000000 -0.041119 -0.078974 -0.198258  0.027200
cp      -0.071966 -0.041119  1.000000  0.038177 -0.081641  0.079294
trestbps 0.271121 -0.078974  0.038177  1.000000  0.127977  0.181767
chol     0.219823 -0.198258 -0.081641  0.127977  1.000000  0.026917
fbs      0.121243  0.027200  0.079294  0.181767  0.026917  1.000000
restecg -0.132696 -0.055117  0.043581 -0.123794 -0.147410 -0.104051
thalach -0.390227 -0.049365  0.306839 -0.039264 -0.021772 -0.008866
exang    0.088163  0.139157 -0.401513  0.061197  0.067382  0.049261
oldpeak  0.208137  0.084687 -0.174733  0.187434  0.064880  0.010859
slope   -0.169105 -0.026666  0.131633 -0.120445 -0.014248 -0.061902
ca       0.271551  0.111729 -0.176206  0.104554  0.074259  0.137156
thal     0.072297  0.198424 -0.163341  0.059276  0.100244 -0.042177
target  -0.229324 -0.279501  0.434854 -0.138772 -0.099966 -0.041164

      restecg  thalach  exang  oldpeak  slope  ca  \
age     -0.132696 -0.390227  0.088163  0.208137 -0.169105  0.271551
sex     -0.055117 -0.049365  0.139157  0.084687 -0.026666  0.111729
cp       0.043581  0.306839 -0.401513 -0.174733  0.131633 -0.176206
trestbps -0.123794 -0.039264  0.061197  0.187434 -0.120445  0.104554
chol     -0.147410 -0.021772  0.067382  0.064880 -0.014248  0.074259
fbs      -0.104051 -0.008866  0.049261  0.010859 -0.061902  0.137156
restecg   1.000000  0.048411 -0.065606 -0.050114  0.086086 -0.078072
thalach   0.048411  1.000000 -0.380281 -0.349796  0.395308 -0.207888
exang    -0.065606 -0.380281  1.000000  0.310844 -0.267335  0.107849
oldpeak  -0.050114 -0.349796  0.310844  1.000000 -0.575189  0.221816
slope     0.086086  0.395308 -0.267335 -0.575189  1.000000 -0.073440
ca       -0.078072 -0.207888  0.107849  0.221816 -0.073440  1.000000
thal     -0.020504 -0.098068  0.197201  0.202672 -0.094090  0.149014
target    0.134468  0.422895 -0.438029 -0.438441  0.345512 -0.382085

      thal  target
age      0.072297 -0.229324
sex      0.198424 -0.279501
cp      -0.163341  0.434854

```

```

trestbps  0.059276 -0.138772
chol       0.100244 -0.099966
fbs        -0.042177 -0.041164
restecg    -0.020504  0.134468
thalach    -0.098068  0.422895
exang      0.197201 -0.438029
oldpeak    0.202672 -0.438441
slope      -0.094090  0.345512
ca         0.149014 -0.382085
thal       1.000000 -0.337838
target     -0.337838  1.000000

```

```

[12]: df['target'].unique()                                     # this is Numerical
      ↪ coloumn bcoz its, countable

```

```

[12]: array([0, 1], dtype=int64)

```

```

[13]: df['oldpeak'].unique()

```

```

[13]: array([1. , 3.1, 2.6, 0. , 1.9, 4.4, 0.8, 3.2, 1.6, 3. , 0.7, 4.2, 1.5,
            2.2, 1.1, 0.3, 0.4, 0.6, 3.4, 2.8, 1.2, 2.9, 3.6, 1.4, 0.2, 2. ,
            5.6, 0.9, 1.8, 6.2, 4. , 2.5, 0.5, 0.1, 2.1, 2.4, 3.8, 2.3, 1.3,
            3.5])

```

```

[14]: df['age'].unique()

```

```

[14]: array([52, 53, 70, 61, 62, 58, 55, 46, 54, 71, 43, 34, 51, 50, 60, 67, 45,
            63, 42, 44, 56, 57, 59, 64, 65, 41, 66, 38, 49, 48, 29, 37, 47, 68,
            76, 40, 39, 77, 69, 35, 74], dtype=int64)

```

```

[15]: df['sex'].unique()

```

```

[15]: array([1, 0], dtype=int64)

```

```

[16]: df['cp'].unique()

```

```

[16]: array([0, 1, 2, 3], dtype=int64)

```

```

[17]: df['trestbps'].unique()

```

```

[17]: array([125, 140, 145, 148, 138, 100, 114, 160, 120, 122, 112, 132, 118,
            128, 124, 106, 104, 135, 130, 136, 180, 129, 150, 178, 146, 117,
            152, 154, 170, 134, 174, 144, 108, 123, 110, 142, 126, 192, 115,
            94, 200, 165, 102, 105, 155, 172, 164, 156, 101], dtype=int64)

```

```

[18]: df['chol'].unique()

```

```
[18]: array([212, 203, 174, 294, 248, 318, 289, 249, 286, 149, 341, 210, 298,
          204, 308, 266, 244, 211, 185, 223, 208, 252, 209, 307, 233, 319,
          256, 327, 169, 131, 269, 196, 231, 213, 271, 263, 229, 360, 258,
          330, 342, 226, 228, 278, 230, 283, 241, 175, 188, 217, 193, 245,
          232, 299, 288, 197, 315, 215, 164, 326, 207, 177, 257, 255, 187,
          201, 220, 268, 267, 236, 303, 282, 126, 309, 186, 275, 281, 206,
          335, 218, 254, 295, 417, 260, 240, 302, 192, 225, 325, 235, 274,
          234, 182, 167, 172, 321, 300, 199, 564, 157, 304, 222, 184, 354,
          160, 247, 239, 246, 409, 293, 180, 250, 221, 200, 227, 243, 311,
          261, 242, 205, 306, 219, 353, 198, 394, 183, 237, 224, 265, 313,
          340, 259, 270, 216, 264, 276, 322, 214, 273, 253, 176, 284, 305,
          168, 407, 290, 277, 262, 195, 166, 178, 141], dtype=int64)
```

```
[19]: df['fbs'].unique()
```

```
[19]: array([0, 1], dtype=int64)
```

```
[20]: df['restecg'].unique()
```

```
[20]: array([1, 0, 2], dtype=int64)
```

```
[21]: df['thalach'].unique()
```

```
[21]: array([168, 155, 125, 161, 106, 122, 140, 145, 144, 116, 136, 192, 156,
          142, 109, 162, 165, 148, 172, 173, 146, 179, 152, 117, 115, 112,
          163, 147, 182, 105, 150, 151, 169, 166, 178, 132, 160, 123, 139,
          111, 180, 164, 202, 157, 159, 170, 138, 175, 158, 126, 143, 141,
          167, 95, 190, 118, 103, 181, 108, 177, 134, 120, 171, 149, 154,
          153, 88, 174, 114, 195, 133, 96, 124, 131, 185, 194, 128, 127,
          186, 184, 188, 130, 71, 137, 99, 121, 187, 97, 90, 129, 113],
          dtype=int64)
```

```
[22]: df['exang'].unique()
```

```
[22]: array([0, 1], dtype=int64)
```

```
[23]: df['slope'].unique()
```

```
[23]: array([2, 0, 1], dtype=int64)
```

```
[24]: df['ca'].unique()
```

```
[24]: array([2, 0, 1, 3, 4], dtype=int64)
```

```
[25]: df['thal'].unique()
```

```
[25]: array([3, 2, 1, 0], dtype=int64)
```

```
[26]: numerics = ['int16', 'int32', 'int64', 'float16', 'float32', 'float64']
newdf = df.select_dtypes(include=numerics)
newdf
```

```
[26]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
0	52	1	0	125	212	0	1	168	0	1.0	
1	53	1	0	140	203	1	0	155	1	3.1	
2	70	1	0	145	174	0	1	125	1	2.6	
3	61	1	0	148	203	0	1	161	0	0.0	
4	62	0	0	138	294	1	1	106	0	1.9	
...	...	...	...	...	...	...	...	...	...	...	...
1020	59	1	1	140	221	0	1	164	1	0.0	
1021	60	1	0	125	258	0	0	141	1	2.8	
1022	47	1	0	110	275	0	0	118	1	1.0	
1023	50	0	0	110	254	0	0	159	0	0.0	
1024	54	1	0	120	188	0	1	113	0	1.4	

	slope	ca	thal	target
0	2	2	3	0
1	0	0	3	0
2	0	0	3	0
3	2	1	3	0
4	1	3	2	0
...	...	...	...	...
1020	2	0	2	1
1021	1	1	3	0
1022	1	1	2	0
1023	2	0	2	1
1024	1	1	3	0

[1025 rows x 14 columns]

### 1.1.2 Data Visualization

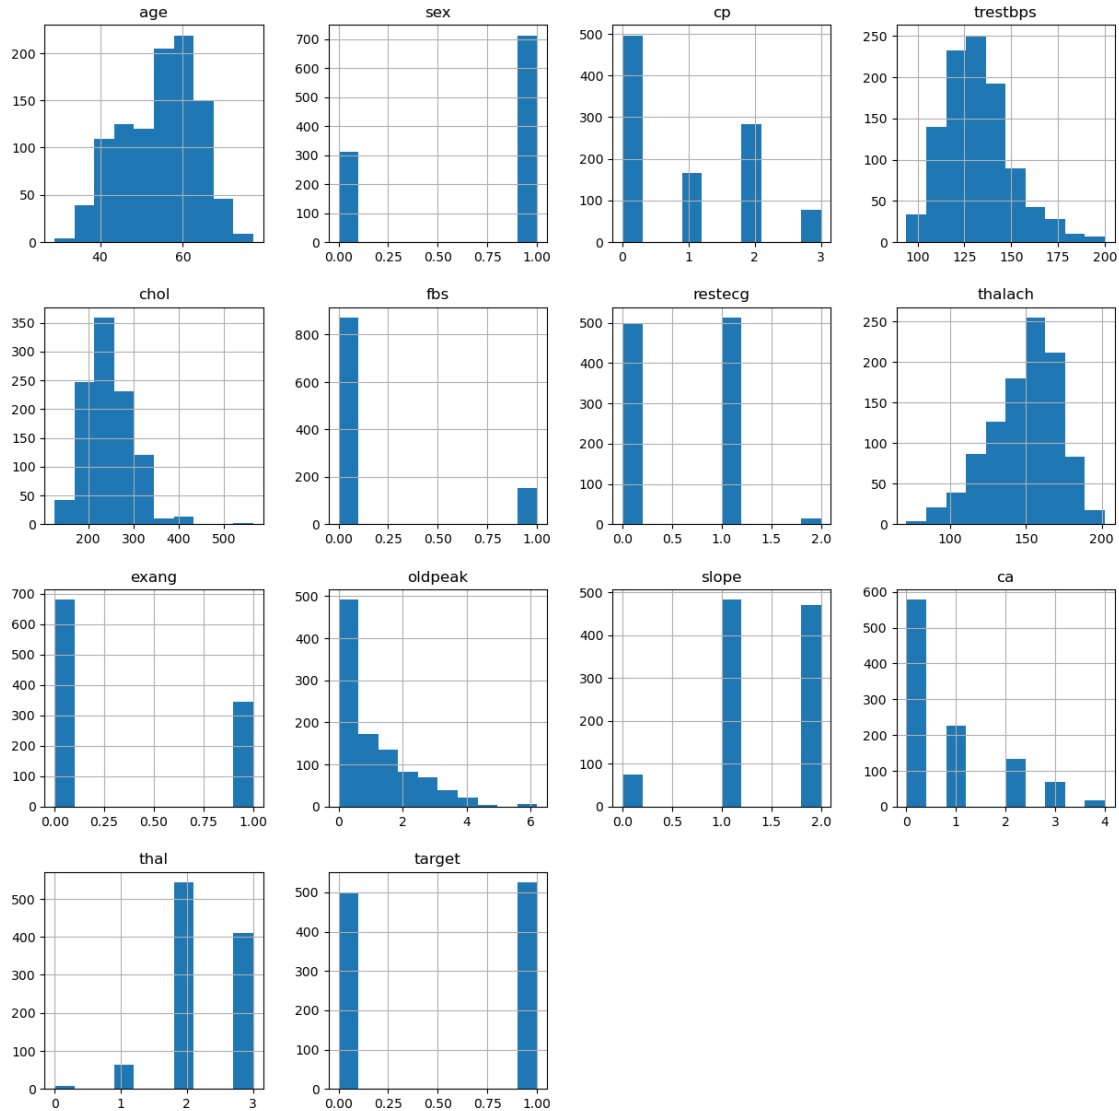
```
[27]: # Importing essential libraries
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
```

```
[28]: fig = plt.figure(figsize = (15,15))
ax = fig.gca()
g = df.hist(ax=ax)
```

C:\Users\DELL\AppData\Local\Temp\ipykernel\_10700\3980286831.py:3: UserWarning:  
To output multiple subplots, the figure containing the passed axes is being  
cleared.

```
g = df.hist(ax=ax)
```





find=iqr #using quantile for removing and detectioning those outliers

```
[29]: r=df.select_dtypes(exclude=['object'])
      r
```

```
[29]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
0	52	1	0	125	212	0	1	168	0	1.0	
1	53	1	0	140	203	1	0	155	1	3.1	
2	70	1	0	145	174	0	1	125	1	2.6	
3	61	1	0	148	203	0	1	161	0	0.0	
4	62	0	0	138	294	1	1	106	0	1.9	
...	...	...	...	...	...	...	...	...	...	...	...
1020	59	1	1	140	221	0	1	164	1	0.0	


1021	60	1	0	125	258	0	0	141	1	2.8
1022	47	1	0	110	275	0	0	118	1	1.0
1023	50	0	0	110	254	0	0	159	0	0.0
1024	54	1	0	120	188	0	1	113	0	1.4

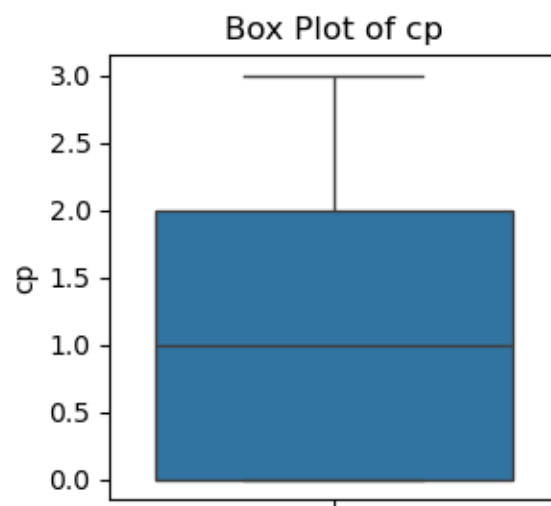

	slope	ca	thal	target
0	2	2	3	0
1	0	0	3	0
2	0	0	3	0
3	2	1	3	0
4	1	3	2	0
...	...	...	...	...
1020	2	0	2	1
1021	1	1	3	0
1022	1	1	2	0
1023	2	0	2	1
1024	1	1	3	0

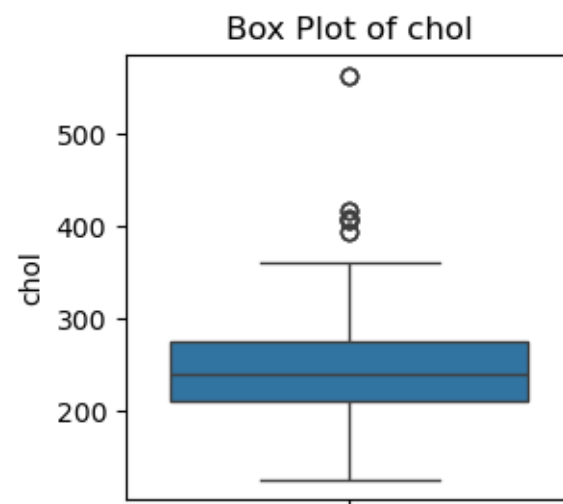
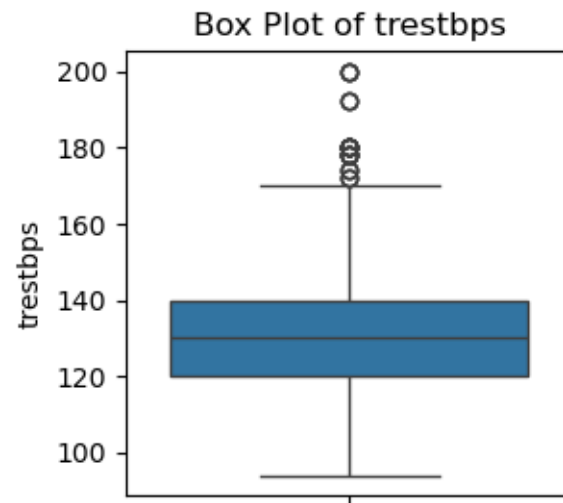
[1025 rows x 14 columns]

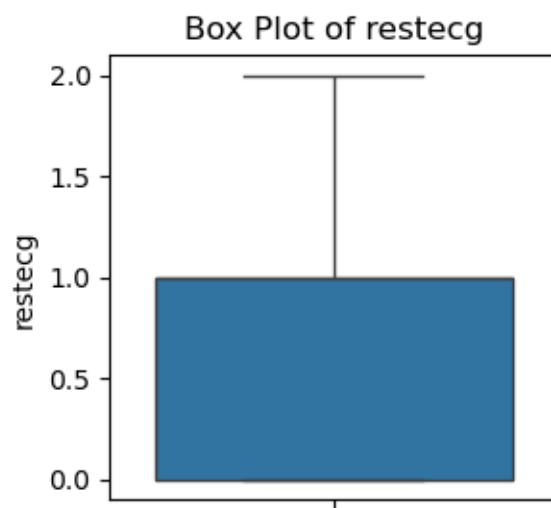
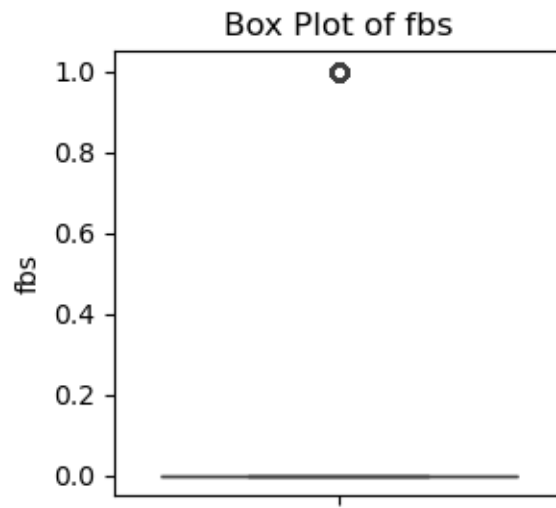
#### outlayer detection

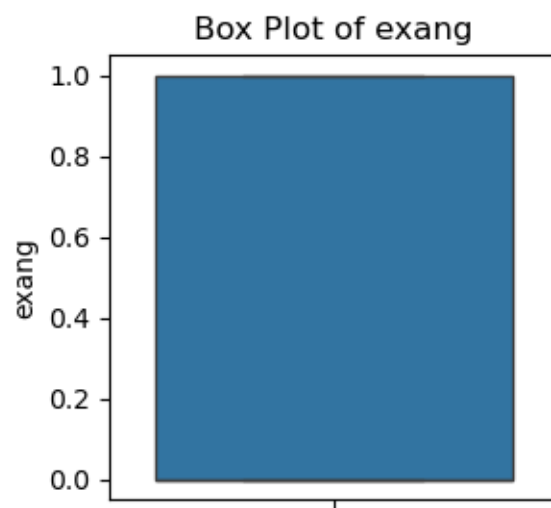
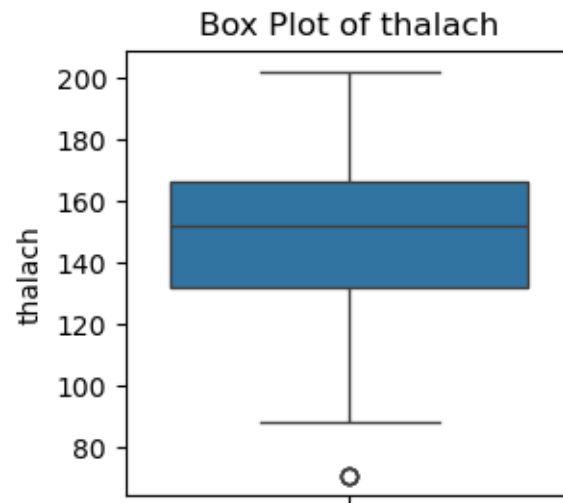
```
[30]: import seaborn as sns
      #outlier detection
      for col in df:
          plt.figure(figsize=(3, 3))
          sns.boxplot(y=df[col])
          plt.title(f'Box Plot of {col}')
          plt.show()
```

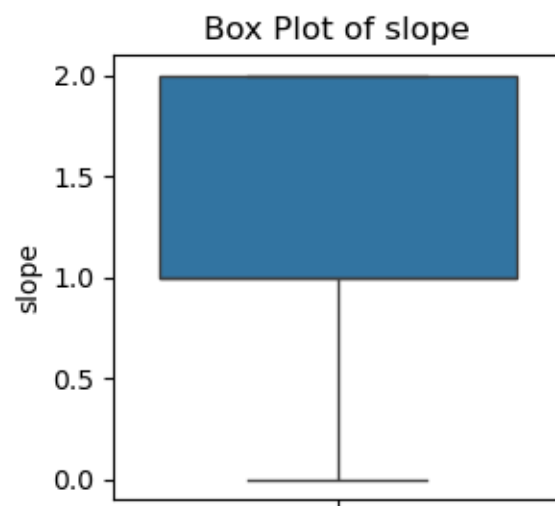
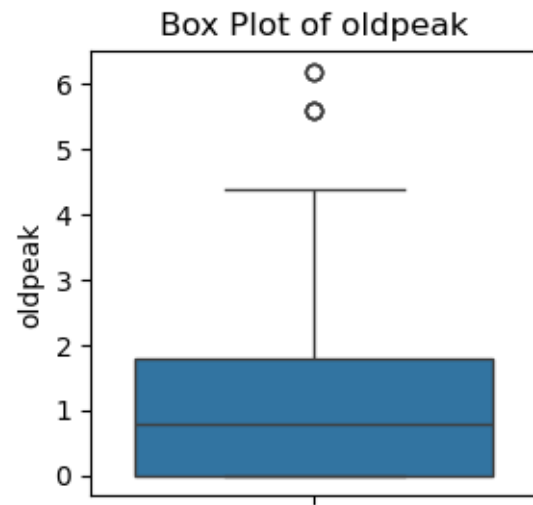


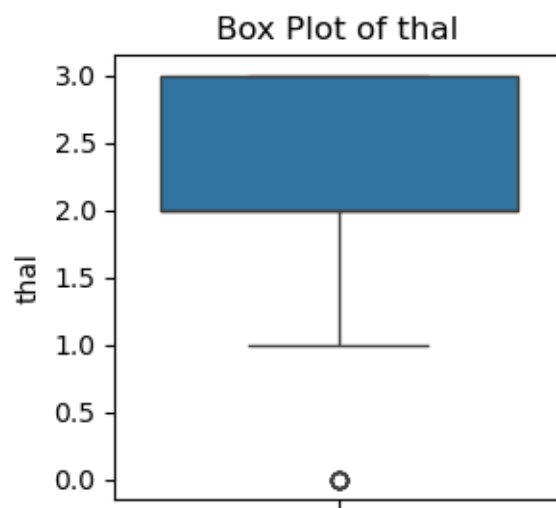
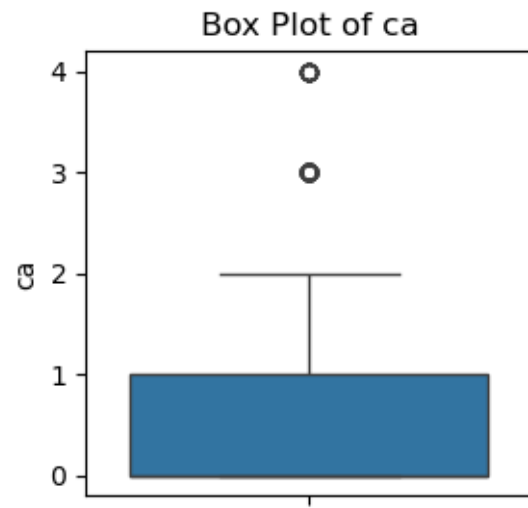




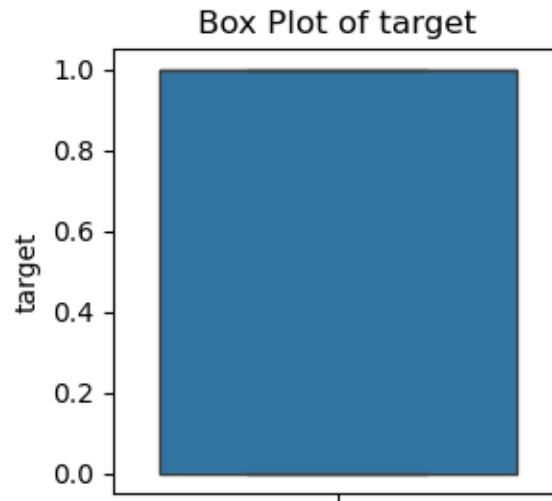












```
[31]: q1=r.quantile(0.25)                                # outlier detection
      q1
```

```
[31]: age          48.0
      sex           0.0
      cp           0.0
      trestbps     120.0
      chol        211.0
      fbs          0.0
      restecg      0.0
      thalach      132.0
      exang        0.0
      oldpeak      0.0
      slope        1.0
      ca           0.0
      thal         2.0
      target       0.0
      Name: 0.25, dtype: float64
```

```
[32]: q3=r.quantile(0.75)
      q3
```

```
[32]: age          61.0
      sex           1.0
      cp           2.0
      trestbps     140.0
      chol        275.0
      fbs          0.0
      restecg      1.0
```

```

thalach      166.0
exang        1.0
oldpeak      1.8
slope        2.0
ca           1.0
thal         3.0
target       1.0
Name: 0.75, dtype: float64

```

```
[33]: IQR=q3-q1
      IQR
```

```
[33]: age      13.0
      sex      1.0
      cp       2.0
      trestbps 20.0
      chol     64.0
      fbs      0.0
      restecg  1.0
      thalach  34.0
      exang    1.0
      oldpeak  1.8
      slope    1.0
      ca       1.0
      thal     1.0
      target   1.0
      dtype: float64

```

```
[34]: a=((r<q1-1.5*IQR)|(r>q1+1.5*IQR))
      a
```

```
[34]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	\
0	False	False	False	False	False	False	False	False	False	
1	False	False	False	False	False	True	False	False	False	
2	True	False	False	False	False	False	False	False	False	
3	False	False	False	False	False	False	False	False	False	
4	False	False	False	False	False	True	False	False	False	
...	...	...	...	...	...	...	...	...	...	
1020	False	False	False	False	False	False	False	False	False	
1021	False	False	False	False	False	False	False	False	False	
1022	False	False	False	False	False	False	False	False	False	
1023	False	False	False	False	False	False	False	False	False	
1024	False	False	False	False	False	False	False	False	False	

	oldpeak	slope	ca	thal	target
0	False	False	True	False	False
1	True	False	False	False	False

2	False	False	False	False	False
3	False	False	False	False	False
4	False	False	True	False	False
...	...	...	...	...	...
1020	False	False	False	False	False
1021	True	False	False	False	False
1022	False	False	False	False	False
1023	False	False	False	False	False
1024	False	False	False	False	False

[1025 rows x 14 columns]

```
[35]: df1=df[~((r<q1-1.5*IQR)|(r>q1+1.5*IQR)).any(axis=1)]
df1
```

```
[35]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
3	61	1	0	148	203	0	1	161	0	0.0	
5	58	0	0	100	248	0	0	122	0	1.0	
8	46	1	0	120	249	0	0	144	0	0.8	
17	54	1	0	124	266	0	0	109	1	2.2	
18	50	0	1	120	244	0	1	162	0	1.1	
...	...	...	...	...	...	...	...	...	...	...	...
1019	47	1	0	112	204	0	1	143	0	0.1	
1020	59	1	1	140	221	0	1	164	1	0.0	
1022	47	1	0	110	275	0	0	118	1	1.0	
1023	50	0	0	110	254	0	0	159	0	0.0	
1024	54	1	0	120	188	0	1	113	0	1.4	

	slope	ca	thal	target
3	2	1	3	0
5	1	0	2	1
8	2	0	3	0
17	1	1	3	0
18	2	0	2	1
...	...	...	...	...
1019	2	0	2	1
1020	2	0	2	1
1022	1	1	2	0
1023	2	0	2	1
1024	1	1	3	0

[474 rows x 14 columns]

```
[36]: print(f"Original data shape: {df.shape}")
print(f"Cleaned data shape: {df1.shape}")
```

Original data shape: (1025, 14)

Cleaned data shape: (474, 14)

## 2 univariant analysis

```
[37]: df1.groupby(['target']).count()
```

```
[37]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
target											
0	173	173	173	173	173	173	173	173	173	173	
1	301	301	301	301	301	301	301	301	301	301	

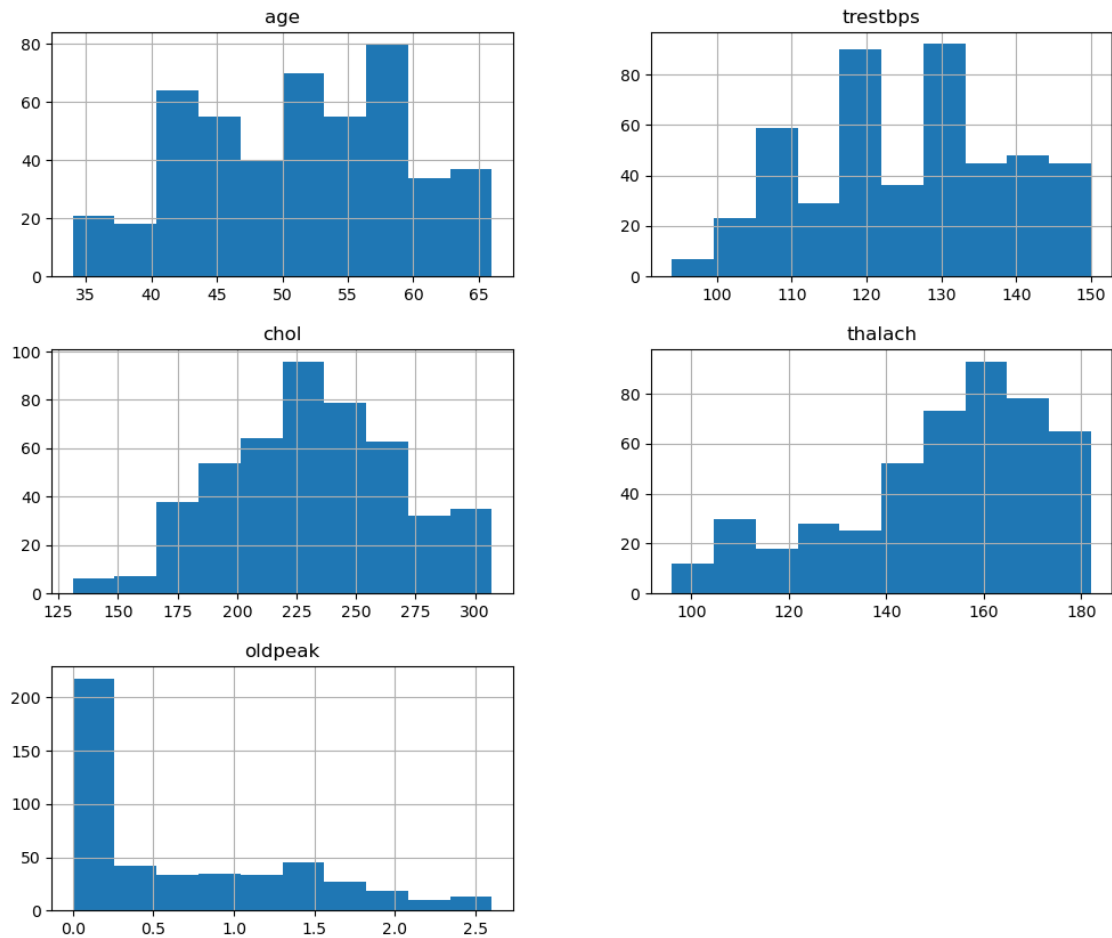
	slope	ca	thal
target			
0	173	173	173
1	301	301	301

```
[38]: a=df1.groupby(['target']).size().reset_index(name='count')
a
```

```
[38]:
```

	target	count
0	0	173
1	1	301

```
[39]: df1[['age', 'trestbps', 'chol', 'thalach', 'oldpeak']].hist(figsize=(12, 10))
plt.show()
```



```
[40]: x=filter['target']
plt.hist(x,bins=2,color="skyblue")
plt.title("Histogram --- clarity of target")
plt.xlabel("target")
plt.ylabel("count")
plt.show()
```

*# Histogram*

```
-----
TypeError                                Traceback (most recent call last)
Cell In[40], line 1
----> 1 x=filter['target']
      2 plt.hist(x,bins=2,color="skyblue")
      3 plt.title("Histogram --- clarity of target")

TypeError: type 'filter' is not subscriptable
```

### 3 Bivariate Analysis

```
[ ]: p=sns.pairplot(df1)
```

```
[ ]: plt.figure(figsize=(10,10))
sns.heatmap(df1.corr(),annot=True,cmap='coolwarm')
plt.title('Correlation Matrix Heatmap')
plt.show()
```

### 4 M.L Model building

```
[41]: x=df.iloc[:, :-1]
x
```

```
[41]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
0	52	1	0	125	212	0	1	168	0	1.0	
1	53	1	0	140	203	1	0	155	1	3.1	
2	70	1	0	145	174	0	1	125	1	2.6	
3	61	1	0	148	203	0	1	161	0	0.0	
4	62	0	0	138	294	1	1	106	0	1.9	
...	...	...	...	...	...	...	...	...	...	...	
1020	59	1	1	140	221	0	1	164	1	0.0	
1021	60	1	0	125	258	0	0	141	1	2.8	
1022	47	1	0	110	275	0	0	118	1	1.0	
1023	50	0	0	110	254	0	0	159	0	0.0	
1024	54	1	0	120	188	0	1	113	0	1.4	

	slope	ca	thal
0	2	2	3
1	0	0	3
2	0	0	3
3	2	1	3
4	1	3	2
...	...	...	...
1020	2	0	2
1021	1	1	3
1022	1	1	2
1023	2	0	2
1024	1	1	3

[1025 rows x 13 columns]

```
[42]: y=df.iloc[:, -1]
y
```

```
[42]: 0      0
      1      0
      2      0
      3      0
      4      0
      ..
      1020    1
      1021    0
      1022    0
      1023    1
      1024    0
      Name: target, Length: 1025, dtype: int64
```

```
[43]: from sklearn.model_selection import train_test_split
```

```
[44]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.
      ↪2,random_state=42)
```

```
[45]: x_train.shape
```

```
[45]: (820, 13)
```

```
[46]: x_train
```

```
[46]:      age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  \
835   49   1   2      118   149   0         0      126     0       0.8
137   64   0   0      180   325   0         1      154     1       0.0
534   54   0   2      108   267   0         0      167     0       0.0
495   59   1   0      135   234   0         1      161     0       0.5
244   51   1   2      125   245   1         0      166     0       2.4
..    ...  ...  ..      ...   ...   ...      ...      ...     ...
700   41   1   2      130   214   0         0      168     0       2.0
71    61   1   0      140   207   0         0      138     1       1.9
106   51   1   0      140   299   0         1      173     1       1.6
270   43   1   0      110   211   0         1      161     0       0.0
860   52   1   0      112   230   0         1      160     0       0.0

      slope  ca  thal
835      2   3    2
137      2   0    2
534      2   0    2
495      1   0    3
244      1   0    2
..      ...  ...  ...
700      1   0    2
71      2   1    3
106      2   0    3
```

```

270      2  0   3
860      2  1   2

```

```
[820 rows x 13 columns]
```

```
[47]: x_test.shape
```

```
[47]: (205, 13)
```

```
[48]: x_test
```

```
[48]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
527	62	0	0	124	209	0	1	163	0	0.0	
359	53	0	2	128	216	0	0	115	0	0.0	
447	55	1	0	160	289	0	0	145	1	0.8	
31	50	0	1	120	244	0	1	162	0	1.1	
621	48	1	0	130	256	1	0	150	1	0.0	
..	...	...	..	...	...	...	...	...			
832	68	1	2	118	277	0	1	151	0	1.0	
796	41	1	1	135	203	0	1	132	0	0.0	
644	44	1	2	120	226	0	1	169	0	0.0	
404	61	1	0	140	207	0	0	138	1	1.9	
842	58	1	2	112	230	0	0	165	0	2.5	

	slope	ca	thal
527	2	0	2
359	2	0	0
447	1	1	3
31	2	0	2
621	2	2	3
..	...	..	...
832	2	1	3
796	1	0	1
644	2	0	2
404	2	1	3
842	1	1	3

```
[205 rows x 13 columns]
```

```
[49]: y_train.shape
```

```
[49]: (820,)
```

```
[50]: y_train
```

```
[50]: 835    0
      137    1
```



```

534    1
495    1
244    1
..
700    1
71     0
106    0
270    1
860    0
Name: target, Length: 820, dtype: int64

```

```
[51]: y_test.shape
```

```
[51]: (205,)
```

```
[52]: y_test
```

```

[52]: 527    1
      359    1
      447    0
      31     1
      621    0
      ..
      832    1
      796    1
      644    1
      404    0
      842    0
Name: target, Length: 205, dtype: int64

```

```
[ ]:
```

## 5 LogisticRegression

```

[53]: from sklearn.model_selection import train_test_split

      # Define your feature columns and target column
      x = df.drop('target', axis=1)
      y = df['target']

      # Split the data into training and testing sets
      x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
      ↪random_state=42)

```

```

[54]: from sklearn.linear_model import LogisticRegression
      from sklearn.metrics import accuracy_score, classification_report

```

```

# Initialize and train the model
logistic_model = LogisticRegression()
logistic_model.fit(x_train, y_train)

# Make predictions
y_pred_logistic = logistic_model.predict(x_test)

# Evaluate the model
accuracy_logistic = accuracy_score(y_test, y_pred_logistic)
report_logistic = classification_report(y_test, y_pred_logistic)

print("Logistic Regression Accuracy:", accuracy_logistic)
print("Logistic Regression Classification Report:\n", report_logistic)

```

Logistic Regression Accuracy: 0.7853658536585366

Logistic Regression Classification Report:

	precision	recall	f1-score	support
0	0.85	0.70	0.76	102
1	0.74	0.87	0.80	103
accuracy			0.79	205
macro avg	0.79	0.78	0.78	205
weighted avg	0.79	0.79	0.78	205

C:\Users\DELL\anaconda3\Lib\site-packages\sklearn\linear\_model\\_logistic.py:458:

ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

## 5.1 WITH HP

```
[55]: from sklearn.model_selection import GridSearchCV
```

```

# Define parameter grid
param_grid_logistic = {
    'C': [0.01, 0.1, 1, 10, 100],
    'penalty': ['l1', 'l2']
}

```

```

# Initialize and train the model with GridSearch
grid_search_logistic = GridSearchCV(LogisticRegression(max_iter=1000),
    ↪param_grid_logistic, cv=5, n_jobs=-1)
grid_search_logistic.fit(x_train, y_train)

# Best parameters and best score
print("Best Parameters for Logistic Regression:", grid_search_logistic.
    ↪best_params_)
print("Best Score for Logistic Regression:", grid_search_logistic.best_score_)

# Make predictions
y_pred_logistic_tuned = grid_search_logistic.predict(x_test)

# Evaluate the model
accuracy_logistic_tuned = accuracy_score(y_test, y_pred_logistic_tuned)
report_logistic_tuned = classification_report(y_test, y_pred_logistic_tuned)

print("Logistic Regression Accuracy (With Tuning):", accuracy_logistic_tuned)
print("Logistic Regression Classification Report (With Tuning):\n",
    ↪report_logistic_tuned)

```

C:\Users\DELL\anaconda3\Lib\site-packages\sklearn\model\_selection\\_validation.py:378: FitFailedWarning:  
 25 fits failed out of a total of 50.  
 The score on these train-test partitions for these parameters will be set to nan.  
 If these failures are not expected, you can try to debug them by setting error\_score='raise'.

Below are more details about the failures:

```

-----
25 fits failed with the following error:
Traceback (most recent call last):
  File "C:\Users\DELL\anaconda3\Lib\site-packages\sklearn\model_selection\_validation.py", line 686, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\DELL\anaconda3\Lib\site-packages\sklearn\linear_model\_logistic.py", line 1162, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
    ~~~~~
  File "C:\Users\DELL\anaconda3\Lib\site-packages\sklearn\linear_model\_logistic.py", line 54, in _check_solver
    raise ValueError(
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

warnings.warn(some_fits_failed_message, FitFailedWarning)
C:\Users\DELL\anaconda3\Lib\site-

```

```
packages\sklearn\model_selection\_search.py:952: UserWarning: One or more of the
test scores are non-finite: [          nan 0.81585366          nan 0.84512195
nan 0.85
```

```
          nan 0.84878049          nan 0.84878049]
warnings.warn(
```

```
Best Parameters for Logistic Regression: {'C': 1, 'penalty': 'l2'}
```

```
Best Score for Logistic Regression: 0.85
```

```
Logistic Regression Accuracy (With Tuning): 0.7951219512195122
```

```
Logistic Regression Classification Report (With Tuning):
```

	precision	recall	f1-score	support
0	0.85	0.72	0.78	102
1	0.76	0.87	0.81	103
accuracy			0.80	205
macro avg	0.80	0.79	0.79	205
weighted avg	0.80	0.80	0.79	205

```
[ ]:
```

## 6 RandomForestClassifier

```
[56]: from sklearn.ensemble import RandomForestClassifier

# Initialize and train the model
rf_model = RandomForestClassifier()
rf_model.fit(x_train, y_train)

# Make predictions
y_pred_rf = rf_model.predict(x_test)

# Evaluate the model
accuracy_rf = accuracy_score(y_test, y_pred_rf)
report_rf = classification_report(y_test, y_pred_rf)

print("Random Forest Accuracy:", accuracy_rf)
print("Random Forest Classification Report:\n", report_rf)
```

```
Random Forest Accuracy: 0.9853658536585366
```

```
Random Forest Classification Report:
```

	precision	recall	f1-score	support
0	0.97	1.00	0.99	102
1	1.00	0.97	0.99	103

accuracy			0.99	205
macro avg	0.99	0.99	0.99	205
weighted avg	0.99	0.99	0.99	205

## 6.1 WITH HP

```
[57]: from sklearn.model_selection import GridSearchCV

# Define parameter grid
param_grid_rf = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5],
    'min_samples_leaf': [1, 2]
}

# Initialize and train the model with GridSearch
grid_search_rf = GridSearchCV(RandomForestClassifier(), param_grid_rf, cv=5,
                               n_jobs=-1)
grid_search_rf.fit(x_train, y_train)

# Best parameters and best score
print("Best Parameters for Random Forest:", grid_search_rf.best_params_)
print("Best Score for Random Forest:", grid_search_rf.best_score_)

# Make predictions
y_pred_rf_tuned = grid_search_rf.predict(x_test)

# Evaluate the model
accuracy_rf_tuned = accuracy_score(y_test, y_pred_rf_tuned)
report_rf_tuned = classification_report(y_test, y_pred_rf_tuned)

print("Random Forest Accuracy (With Tuning):", accuracy_rf_tuned)
print("Random Forest Classification Report (With Tuning):\n", report_rf_tuned)
```

Best Parameters for Random Forest: {'max\_depth': 10, 'min\_samples\_leaf': 1, 'min\_samples\_split': 2, 'n\_estimators': 100}

Best Score for Random Forest: 0.9829268292682928

Random Forest Accuracy (With Tuning): 0.9853658536585366

Random Forest Classification Report (With Tuning):

	precision	recall	f1-score	support
0	0.97	1.00	0.99	102
1	1.00	0.97	0.99	103
accuracy			0.99	205
macro avg	0.99	0.99	0.99	205

weighted avg	0.99	0.99	0.99	205
--------------	------	------	------	-----

[ ]:

## 7 SVC Classifier

```
[58]: from sklearn.svm import SVC

# Initialize and train the model
svm_model = SVC()
svm_model.fit(x_train, y_train)

# Make predictions
y_pred_svm = svm_model.predict(x_test)

# Evaluate the model
accuracy_svm = accuracy_score(y_test, y_pred_svm)
report_svm = classification_report(y_test, y_pred_svm)

print("SVM Accuracy:", accuracy_svm)
print("SVM Classification Report:\n", report_svm)
```

SVM Accuracy: 0.6829268292682927

SVM Classification Report:

	precision	recall	f1-score	support
0	0.71	0.61	0.66	102
1	0.66	0.76	0.71	103
accuracy			0.68	205
macro avg	0.69	0.68	0.68	205
weighted avg	0.69	0.68	0.68	205

### 7.1 WITH HP

```
[59]: from sklearn.model_selection import GridSearchCV

# Define parameter grid
param_grid_svm = {
    'C': [0.1, 1, 10],
    'kernel': ['linear'],
}

# Initialize and train the model with GridSearch
```

```

grid_search_svm = GridSearchCV(SVC(), param_grid_svm, cv=5, n_jobs=-1)
grid_search_svm.fit(x_train, y_train)

# Best parameters and best score
print("Best Parameters for SVM:", grid_search_svm.best_params_)
print("Best Score for SVM:", grid_search_svm.best_score_)

# Make predictions
y_pred_svm_tuned = grid_search_svm.predict(x_test)

# Evaluate the model
accuracy_svm_tuned = accuracy_score(y_test, y_pred_svm_tuned)
report_svm_tuned = classification_report(y_test, y_pred_svm_tuned)

print("SVM Accuracy (With Tuning):", accuracy_svm_tuned)
print("SVM Classification Report (With Tuning):\n", report_svm_tuned)

```

```

Best Parameters for SVM: {'C': 0.1, 'kernel': 'linear'}
Best Score for SVM: 0.8512195121951219
SVM Accuracy (With Tuning): 0.7951219512195122
SVM Classification Report (With Tuning):

```

	precision	recall	f1-score	support
0	0.88	0.68	0.77	102
1	0.74	0.91	0.82	103
accuracy			0.80	205
macro avg	0.81	0.79	0.79	205
weighted avg	0.81	0.80	0.79	205

[ ]:

## 8 KNN Classifier

```

[60]: from sklearn.neighbors import KNeighborsClassifier

# Initialize and train the model
knn_model = KNeighborsClassifier()
knn_model.fit(x_train, y_train)

# Make predictions
y_pred_knn = knn_model.predict(x_test)

# Evaluate the model
accuracy_knn = accuracy_score(y_test, y_pred_knn)

```

```
report_knn = classification_report(y_test, y_pred_knn)

print("KNN Accuracy:", accuracy_knn)
print("KNN Classification Report:\n", report_knn)
```

KNN Accuracy: 0.7317073170731707

KNN Classification Report:

	precision	recall	f1-score	support
0	0.73	0.73	0.73	102
1	0.73	0.74	0.73	103
accuracy			0.73	205
macro avg	0.73	0.73	0.73	205
weighted avg	0.73	0.73	0.73	205


## 8.1 With HyperParametric Tuning

```
[61]: error_rate = []
      for i in range(1,40):
          knn = KNeighborsClassifier(n_neighbors=i)
          knn.fit(x_train,y_train)
          pred_i=knn.predict(x_test)
          error_rate.append(np.mean(pred_i !=y_test))
```

```
[62]: plt.figure(figsize=(10,6))
      plt.plot(range(1,40),error_rate,color='blue',linestyle='dashed',marker='o',
                markerfacecolor='red',markersize=10)
      plt.title('Error Rate vs. K Value')
      plt.xlabel('K')
      plt.ylabel('Error Rate')
```

```
[62]: Text(0, 0.5, 'Error Rate')
```





```
[63]: from sklearn.metrics import confusion_matrix
knn= KNeighborsClassifier(n_neighbors=3)

knn.fit(x_train,y_train)
pred = knn.predict(x_test)

print('With K=3')
print('\n')
print(confusion_matrix(y_test,pred))
print(classification_report(y_test,pred))
```

With K=3

```
[[91 11]
 [ 9 94]]
```

	precision	recall	f1-score	support
0	0.91	0.89	0.90	102
1	0.90	0.91	0.90	103
accuracy			0.90	205
macro avg	0.90	0.90	0.90	205
weighted avg	0.90	0.90	0.90	205

[ ]:

[ ]: